

homelle

December 12, 2023

Data Collecting

```
[123]: import numpy as np
import pandas as pd
```

```
[124]: #Reading CASE-SHILLER Index into a dataframe
df_CS = pd.read_csv(r"C:\Users\Prachi\Downloads\CSUSHPISA (1).csv")

#Changing dtype of date column
df_CS["DATE"] = pd.to_datetime(df_CS["DATE"])

#Selecting data post JUNE 2001
mask = df_CS["DATE"] >= "2001-07-01"
df_CS = df_CS[mask]

#Resetting Index
df_CS.reset_index(inplace = True)
df_CS.drop(columns = ["index"], inplace = True)

# Creating "Year" and "Month" columns
df_CS["Year"] = pd.DatetimeIndex(df_CS["DATE"]).year
df_CS["Month"] = pd.DatetimeIndex(df_CS["DATE"]).month
print(df_CS.shape)
df_CS.tail()
```

(247, 4)

```
[124]:
```

	DATE	CSUSHPISA	Year	Month
242	2022-02-01	291.153	2022	2
243	2022-03-01	296.445	2022	3
244	2022-04-01	300.573	2022	4
245	2022-05-01	303.762	2022	5
246	2022-06-01	304.724	2022	6

```
[125]: # Reading Unemployment Rate Data into a dataframe
df_unemp = pd.read_csv(r"C:\Users\Prachi\Downloads\UNRATE (2).csv", names =
    ↪ ["DATE", "UNRATE"], skiprows = 1)
df_unemp.drop(df_unemp.index[267:912], inplace = True)
```

```
print(df_unemp.shape)
df_unemp.tail()
```

(247, 2)

```
[125]:
```

	DATE	UNRATE
242	2022-02-01	3.8
243	2022-03-01	3.6
244	2022-04-01	3.6
245	2022-05-01	3.6
246	2022-06-01	3.6

```
[126]: # Reading Per Capita GDP Data into a dataframe
df_pcgdp = pd.read_csv(r"C:\Users\Prachi\Downloads\A939RX0Q048SBEA.csv", names_
↳= ["DATE", "Per_Capita_GDP"], skiprows = 1)
print(df_pcgdp.shape)
df_pcgdp.tail()
```

(83, 2)

```
[126]:
```

	DATE	Per_Capita_GDP
78	2021-04-01	64157.0
79	2021-07-01	64615.0
80	2021-10-01	65651.0
81	2022-01-01	65286.0
82	2022-04-01	65127.0

```
[127]: # Interest Rate Data
df_Fed_rate = pd.read_csv(r"C:\Users\Prachi\Downloads\FEDFUNDS (1).csv")
df_Fed_rate.drop(df_Fed_rate.index[267:1828], inplace = True)
print(df_Fed_rate.shape)
df_Fed_rate.tail()
```

(247, 2)

```
[127]:
```

	DATE	FEDFUNDS
242	2022-02-01	0.08
243	2022-03-01	0.20
244	2022-04-01	0.33
245	2022-05-01	0.77
246	2022-06-01	1.21

```
[128]: # Reading Per Capita GDP Data into a dataframe
df_cons_price_index = pd.read_csv(r"C:\Users\Prachi\Downloads\WPUSI012011.csv",
↳names = ["DATE", "Cons_Materials"], skiprows = 1)
df_cons_price_index.drop(df_cons_price_index.index[267:611], inplace = True)
print(df_cons_price_index.shape)
```

```
df_cons_price_index.tail()
```

(247, 2)

```
[128]:
```

	DATE	Cons_Materials
242	2022-02-01	343.583
243	2022-03-01	345.852
244	2022-04-01	343.786
245	2022-05-01	353.015
246	2022-06-01	349.800

```
[129]: # Consumer Price Index
df_CPI = pd.read_csv(r"C:\Users\Prachi\Downloads\CPIAUCSL (1).csv", names =
    ↳ ["DATE", "CPI"], skiprows = 1)
df_CPI.drop(df_CPI.index[267:923], inplace = True)
print(df_CPI.shape)
df_CPI.tail()
```

(247, 2)

```
[129]:
```

	DATE	CPI
242	2022-02-01	284.610
243	2022-03-01	287.472
244	2022-04-01	288.611
245	2022-05-01	291.268
246	2022-06-01	294.728

```
[130]: # Monthly new house supply
df_house = pd.read_csv(r"C:\Users\Prachi\Downloads\MSACSR (1).csv", names =
    ↳ ["DATE", "Houses"], skiprows = 1)
df_house.drop(df_house.index[267:731], inplace = True)
print(df_house.shape)
df_house.tail()
```

(247, 2)

```
[130]:
```

	DATE	Houses
242	2022-02-01	6.2
243	2022-03-01	7.0
244	2022-04-01	8.5
245	2022-05-01	8.3
246	2022-06-01	9.5

```
[131]: # Population above 65

df_oldpop = pd.read_csv(r"C:\Users\Prachi\Downloads\SPPPOP65UPTOZSUSA (1).csv",
    ↳ names = ["DATE", "old_percent"], skiprows = 1)
```

```
print(df_oldpop.shape)
df_oldpop.tail()
```

(22, 2)

```
[131]:
```

	DATE	old_percent
17	2018-01-01	15.397698
18	2019-01-01	15.791801
19	2020-01-01	16.223400
20	2021-01-01	16.678895
21	2022-01-01	17.128121

```
[132]: # Urban Population Percent

df_urban = pd.read_csv(r"C:\Users\Prachi\Downloads\LREM64TTUSM156S.csv", names_
    ↪= ["DATE", "Urban_pop"], skiprows = 1)
df_urban.drop(df_urban.index[267:851], inplace = True)
print(df_urban.shape)
df_urban.tail()
```

(247, 2)

```
[132]:
```

	DATE	Urban_pop
242	2022-02-01	70.900503
243	2022-03-01	71.266867
244	2022-04-01	71.222821
245	2022-05-01	71.370775
246	2022-06-01	71.228684

```
[133]: # Housing Subsidies

df_subsidy = pd.read_csv(r"C:\Users\Prachi\Downloads\L312051A027NBEA (1).csv",
    ↪names = ["DATE", "Subsidy"], skiprows = 1)
print(df_subsidy.shape)
df_subsidy.tail()
```

(22, 2)

```
[133]:
```

	DATE	Subsidy
17	2018-01-01	38.859
18	2019-01-01	40.185
19	2020-01-01	44.147
20	2021-01-01	45.299
21	2022-01-01	48.021

```
[134]: # Working age population
```

```
df_working = pd.read_csv(r"C:\Users\Prachi\Downloads\LFWA64TTUSM647S (2).csv",
    ↪names = ["DATE", "Working_Population"], skiprows = 1)
df_working.drop(df_working.index[267:563], inplace = True)
print(df_working.shape)
df_working.tail()
```

(247, 2)

```
[134]:
```

	DATE	Working_Population
242	2022-02-01	2.071042e+08
243	2022-03-01	2.070130e+08
244	2022-04-01	2.070650e+08
245	2022-05-01	2.072705e+08
246	2022-06-01	2.073947e+08

```
[135]: # Real Median Household Income

df_income = pd.read_csv(r"C:\Users\Prachi\Downloads\MEHOINUSA672N.csv", names =
    ↪["DATE", "Income"], skiprows = 1)
print(df_income.shape)
df_income.tail()
```

(22, 2)

```
[135]:
```

	DATE	Income
17	2018-01-01	73030
18	2019-01-01	78250
19	2020-01-01	76660
20	2021-01-01	76330
21	2022-01-01	74580

```
[136]: # Number of households

df_households = pd.read_csv(r"C:\Users\Prachi\Downloads\TTLHH (1).csv", names =
    ↪["DATE", "Num_Households"], skiprows = 1)
print(df_households.shape)
df_households.tail()
```

(22, 2)

```
[136]:
```

	DATE	Num_Households
17	2018-01-01	127586.0
18	2019-01-01	128579.0
19	2020-01-01	128451.0
20	2021-01-01	129224.0
21	2022-01-01	131202.0

Data Preprocessing

```
[137]: # Merging Per Capita GDP (Quarterly data)
df_pcgdp["DATE"] = pd.to_datetime(df_pcgdp["DATE"])
df_CS = pd.merge(df_CS, df_pcgdp, how = "left")
df_CS.head()
```

```
[137]:
```

	DATE	CSUSHPISA	Year	Month	Per_Capita_GDP
0	2001-12-01	116.455	2001	12	NaN
1	2002-01-01	117.144	2002	1	50091.0
2	2002-02-01	117.845	2002	2	NaN
3	2002-03-01	118.687	2002	3	NaN
4	2002-04-01	119.611	2002	4	50286.0

```
[138]: df = pd.DataFrame()
df_bymonth = [df_CS, df_working, df_house, df_CPI, df_unemp,
df_cons_price_index, df_Fed_rate]
for df1 in df_bymonth:
    df1["DATE"] = pd.to_datetime(df1["DATE"])
    df1 = df1.set_index("DATE")
    df = pd.concat([df, df1], axis = 1)
print(df.shape)
df.head(5)
```

(247, 10)

```
[138]:
```

	DATE	CSUSHPISA	Year	Month	Per_Capita_GDP	Working_Population \
	2001-12-01	116.455	2001	12	NaN	1.826419e+08
	2002-01-01	117.144	2002	1	50091.0	1.825664e+08
	2002-02-01	117.845	2002	2	NaN	1.827984e+08
	2002-03-01	118.687	2002	3	NaN	1.830783e+08
	2002-04-01	119.611	2002	4	50286.0	1.832605e+08

	DATE	Houses	CPI	UNRATE	Cons_Materials	FEDFUNDS
	2001-12-01	3.8	177.4	5.7	141.7	1.82
	2002-01-01	4.2	177.7	5.7	142.0	1.73
	2002-02-01	4.0	178.0	5.7	142.2	1.74
	2002-03-01	4.1	178.5	5.7	143.2	1.73
	2002-04-01	4.3	179.3	5.9	143.5	1.75

```
[139]: # Merging other dataframes
others = [df_urban, df_households, df_income, df_subsidy, df_oldpop]
for df1 in others:
    if "Year" not in df1.columns:
        df1["Year"] = pd.DatetimeIndex(df1["DATE"]).year
        df1.set_index("DATE", inplace = True)
        df = pd.merge(df, df1, how = "left", on = "Year")
```

```

else:
    df1.set_index("DATE", inplace = True)
    df = pd.merge(df, df1, how = "left", on = "Year")
df["DATE"] = df_CS["DATE"]
df.set_index("DATE", inplace = True)
df.head()

```

```

[139]:
          CSUSHPISA  Year  Month  Per_Capita_GDP  Working_Population  \
DATE
2001-12-01    116.455  2001     12             NaN        1.826419e+08
2002-01-01    117.144  2002      1        50091.0        1.825664e+08
2002-02-01    117.144  2002      1        50091.0        1.825664e+08
2002-03-01    117.144  2002      1        50091.0        1.825664e+08
2002-04-01    117.144  2002      1        50091.0        1.825664e+08

```

```

          Houses    CPI  UNRATE  Cons_Materials  FEDFUNDS  Urban_pop  \
DATE
2001-12-01     3.8  177.4     5.7           141.7      1.82  72.226108
2002-01-01     4.2  177.7     5.7           142.0      1.73  72.016140
2002-02-01     4.2  177.7     5.7           142.0      1.73  72.333005
2002-03-01     4.2  177.7     5.7           142.0      1.73  72.131883
2002-04-01     4.2  177.7     5.7           142.0      1.73  71.914847

```

```

          Num_Households  Income  Subsidy  old_percent
DATE
2001-12-01        108209.0   66360   20.573   12.296945
2002-01-01        109297.0   65820   24.183   12.287458
2002-02-01        109297.0   65820   24.183   12.287458
2002-03-01        109297.0   65820   24.183   12.287458
2002-04-01        109297.0   65820   24.183   12.287458

```

```

[140]: print(df.shape)

```

```

(2917, 15)

```

```

[141]: df.isna().sum()

```

```

[141]: CSUSHPISA          0
Year                  0
Month                0
Per_Capita_GDP       1945
Working_Population    0
Houses               0
CPI                  0
UNRATE               0
Cons_Materials        0
FEDFUNDS              0

```

```

Urban_pop          0
Num_Households     0
Income             0
Subsidy            0
old_percent        0
dtype: int64

```

The “Per_Capita_GDP” column has missing values because the data was quarterly, We will first fill in the missing values in the “Per_Capita_GDP” column using linear interpolation.

```
[142]: # Filling missing values in the Per_Capita_GDP column using linear interpolation
df["Per_Capita_GDP"] = df["Per_Capita_GDP"].interpolate()
```

```
[143]: df
```

```
[143]:
```

	CSUSHPISA	Year	Month	Per_Capita_GDP	Working_Population	\
DATE						
2001-12-01	116.455	2001	12	NaN	1.826419e+08	
2002-01-01	117.144	2002	1	50091.0	1.825664e+08	
2002-02-01	117.144	2002	1	50091.0	1.825664e+08	
2002-03-01	117.144	2002	1	50091.0	1.825664e+08	
2002-04-01	117.144	2002	1	50091.0	1.825664e+08	
...	
NaT	304.724	2022	6	65127.0	2.073947e+08	
NaT	304.724	2022	6	65127.0	2.073947e+08	
NaT	304.724	2022	6	65127.0	2.073947e+08	
NaT	304.724	2022	6	65127.0	2.073947e+08	
NaT	304.724	2022	6	65127.0	2.073947e+08	

	Houses	CPI	UNRATE	Cons_Materials	FEDFUNDS	Urban_pop	\
DATE							
2001-12-01	3.8	177.400	5.7	141.7	1.82	72.226108	
2002-01-01	4.2	177.700	5.7	142.0	1.73	72.016140	
2002-02-01	4.2	177.700	5.7	142.0	1.73	72.333005	
2002-03-01	4.2	177.700	5.7	142.0	1.73	72.131883	
2002-04-01	4.2	177.700	5.7	142.0	1.73	71.914847	
...		
NaT	9.5	294.728	3.6	349.8	1.21	70.900503	
NaT	9.5	294.728	3.6	349.8	1.21	71.266867	
NaT	9.5	294.728	3.6	349.8	1.21	71.222821	
NaT	9.5	294.728	3.6	349.8	1.21	71.370775	
NaT	9.5	294.728	3.6	349.8	1.21	71.228684	

	Num_Households	Income	Subsidy	old_percent
DATE				
2001-12-01	108209.0	66360	20.573	12.296945
2002-01-01	109297.0	65820	24.183	12.287458

2002-02-01	109297.0	65820	24.183	12.287458
2002-03-01	109297.0	65820	24.183	12.287458
2002-04-01	109297.0	65820	24.183	12.287458
...
NaT	131202.0	74580	48.021	17.128121
NaT	131202.0	74580	48.021	17.128121
NaT	131202.0	74580	48.021	17.128121
NaT	131202.0	74580	48.021	17.128121
NaT	131202.0	74580	48.021	17.128121

[2917 rows x 15 columns]

```
[144]: df.dropna(inplace = True)
```

```
[145]: df.isna().sum()
```

```
[145]: CSUSHPISA      0
Year              0
Month             0
Per_Capita_GDP    0
Working_Population 0
Houses            0
CPI               0
UNRATE            0
Cons_Materials     0
FEDFUNDS          0
Urban_pop         0
Num_Households    0
Income            0
Subsidy           0
old_percent       0
dtype: int64
```

```
[146]: df
```

```
[146]:      CSUSHPISA  Year  Month  Per_Capita_GDP  Working_Population \
DATE
2002-01-01    117.144  2002      1           50091.0      1.825664e+08
2002-02-01    117.144  2002      1           50091.0      1.825664e+08
2002-03-01    117.144  2002      1           50091.0      1.825664e+08
2002-04-01    117.144  2002      1           50091.0      1.825664e+08
2002-05-01    117.144  2002      1           50091.0      1.825664e+08
...          ...   ...   ...           ...           ...
NaT          304.724  2022      6           65127.0      2.073947e+08
NaT          304.724  2022      6           65127.0      2.073947e+08
NaT          304.724  2022      6           65127.0      2.073947e+08
NaT          304.724  2022      6           65127.0      2.073947e+08
```

NaT	304.724	2022	6	65127.0	2.073947e+08	
	Houses	CPI	UNRATE	Cons_Materials	FEDFUNDS	Urban_pop \
DATE						
2002-01-01	4.2	177.700	5.7	142.0	1.73	72.016140
2002-02-01	4.2	177.700	5.7	142.0	1.73	72.333005
2002-03-01	4.2	177.700	5.7	142.0	1.73	72.131883
2002-04-01	4.2	177.700	5.7	142.0	1.73	71.914847
2002-05-01	4.2	177.700	5.7	142.0	1.73	72.026012
...
NaT	9.5	294.728	3.6	349.8	1.21	70.900503
NaT	9.5	294.728	3.6	349.8	1.21	71.266867
NaT	9.5	294.728	3.6	349.8	1.21	71.222821
NaT	9.5	294.728	3.6	349.8	1.21	71.370775
NaT	9.5	294.728	3.6	349.8	1.21	71.228684

	Num_Households	Income	Subsidy	old_percent
DATE				
2002-01-01	109297.0	65820	24.183	12.287458
2002-02-01	109297.0	65820	24.183	12.287458
2002-03-01	109297.0	65820	24.183	12.287458
2002-04-01	109297.0	65820	24.183	12.287458
2002-05-01	109297.0	65820	24.183	12.287458
...
NaT	131202.0	74580	48.021	17.128121
NaT	131202.0	74580	48.021	17.128121
NaT	131202.0	74580	48.021	17.128121
NaT	131202.0	74580	48.021	17.128121
NaT	131202.0	74580	48.021	17.128121

[2916 rows x 15 columns]

```
[147]: df.to_csv("prepared_dataset.csv")
```

```
[148]: us_house_price_df = pd.read_csv("prepared_dataset.csv").set_index("DATE")
us_house_price_df.head()
```

```
[148]: CSUSHPISA  Year  Month  Per_Capita_GDP  Working_Population \
DATE
2002-01-01    117.144  2002      1          50091.0          1.825664e+08
2002-02-01    117.144  2002      1          50091.0          1.825664e+08
2002-03-01    117.144  2002      1          50091.0          1.825664e+08
2002-04-01    117.144  2002      1          50091.0          1.825664e+08
2002-05-01    117.144  2002      1          50091.0          1.825664e+08
```

	Houses	CPI	UNRATE	Cons_Materials	FEDFUNDS	Urban_pop \
DATE						

2002-01-01	4.2	177.7	5.7	142.0	1.73	72.016140
2002-02-01	4.2	177.7	5.7	142.0	1.73	72.333005
2002-03-01	4.2	177.7	5.7	142.0	1.73	72.131883
2002-04-01	4.2	177.7	5.7	142.0	1.73	71.914847
2002-05-01	4.2	177.7	5.7	142.0	1.73	72.026012

	Num_Households	Income	Subsidy	old_percent
DATE				
2002-01-01	109297.0	65820	24.183	12.287458
2002-02-01	109297.0	65820	24.183	12.287458
2002-03-01	109297.0	65820	24.183	12.287458
2002-04-01	109297.0	65820	24.183	12.287458
2002-05-01	109297.0	65820	24.183	12.287458

```
[149]: #Importing necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_squared_error
```

```
[150]: us_house_price_df = pd.read_csv("prepared_dataset.csv").set_index("DATE")
us_house_price_df.head()
```

```
[150]: CSUSHPISA  Year  Month  Per_Capita_GDP  Working_Population  \
DATE
2002-01-01    117.144  2002      1           50091.0        1.825664e+08
2002-02-01    117.144  2002      1           50091.0        1.825664e+08
2002-03-01    117.144  2002      1           50091.0        1.825664e+08
2002-04-01    117.144  2002      1           50091.0        1.825664e+08
2002-05-01    117.144  2002      1           50091.0        1.825664e+08
```

	Houses	CPI	UNRATE	Cons_Materials	FEDFUNDS	Urban_pop	\
DATE							
2002-01-01	4.2	177.7	5.7	142.0	1.73	72.016140	
2002-02-01	4.2	177.7	5.7	142.0	1.73	72.333005	
2002-03-01	4.2	177.7	5.7	142.0	1.73	72.131883	
2002-04-01	4.2	177.7	5.7	142.0	1.73	71.914847	
2002-05-01	4.2	177.7	5.7	142.0	1.73	72.026012	

	Num_Households	Income	Subsidy	old_percent
DATE				
2002-01-01	109297.0	65820	24.183	12.287458
2002-02-01	109297.0	65820	24.183	12.287458
2002-03-01	109297.0	65820	24.183	12.287458
2002-04-01	109297.0	65820	24.183	12.287458

2002-05-01 109297.0 65820 24.183 12.287458

```
[151]: # Dropping year and month columns
us_house_price_df.drop(columns = ["Year", "Month"], inplace = True)
```

Exploratory Data Analysis (EDA) Summary Statistics: Calculate and display summary statistics for each variable, including mean, median, standard deviation, minimum, and maximum values. This gives you an overview of the data's central tendencies and variability.

```
[152]: # Summary statistics
summary_stats = us_house_price_df.describe()
summary_stats
```

```
[152]:
```

	CSUSHPISA	Per_Capita_GDP	Working_Population	Houses	\
count	2916.000000	2916.000000	2.916000e+03	2916.000000	
mean	173.393158	56733.444444	1.990766e+08	5.939300	
std	35.656839	3894.028787	6.902281e+06	1.900553	
min	117.144000	50091.000000	1.825664e+08	3.300000	
25%	146.394000	54100.000000	1.949518e+08	4.500000	
50%	168.634000	55575.500000	2.011083e+08	5.400000	
75%	187.993000	59440.540000	2.053709e+08	6.700000	
max	304.724000	65651.000000	2.073947e+08	12.200000	

	CPI	UNRATE	Cons_Materials	FEDFUNDS	Urban_pop	\
count	2916.000000	2916.000000	2916.000000	2916.000000	2916.000000	
mean	225.173831	6.089300	204.752848	1.290226	69.550604	
std	26.279576	1.963623	39.039207	1.538251	2.079692	
min	177.700000	3.500000	142.000000	0.050000	60.193856	
25%	203.800000	4.700000	182.500000	0.120000	67.699590	
50%	228.329000	5.600000	205.500000	0.650000	70.097404	
75%	244.004000	7.300000	219.700000	1.910000	71.363247	
max	294.728000	14.700000	353.015000	5.260000	72.333005	

	Num_Households	Income	Subsidy	old_percent
count	2916.000000	2916.000000	2916.000000	2916.000000
mean	120384.419753	68561.728395	33.665000	13.828983
std	6278.891569	4444.196095	5.771877	1.436413
min	109297.000000	63350.000000	24.183000	12.277934
25%	116011.000000	65760.000000	29.512000	12.507804
50%	121084.000000	66780.000000	33.283000	13.584437
75%	126224.000000	72090.000000	37.550000	15.066290
max	131202.000000	78250.000000	48.021000	17.128121

Correlation Analysis: Calculate the correlation matrix to measure the linear relationships between variables.

```
[153]: # Correlation matrix
corr_matrix = us_house_price_df.corr()
```

corr_matrix

[153]:

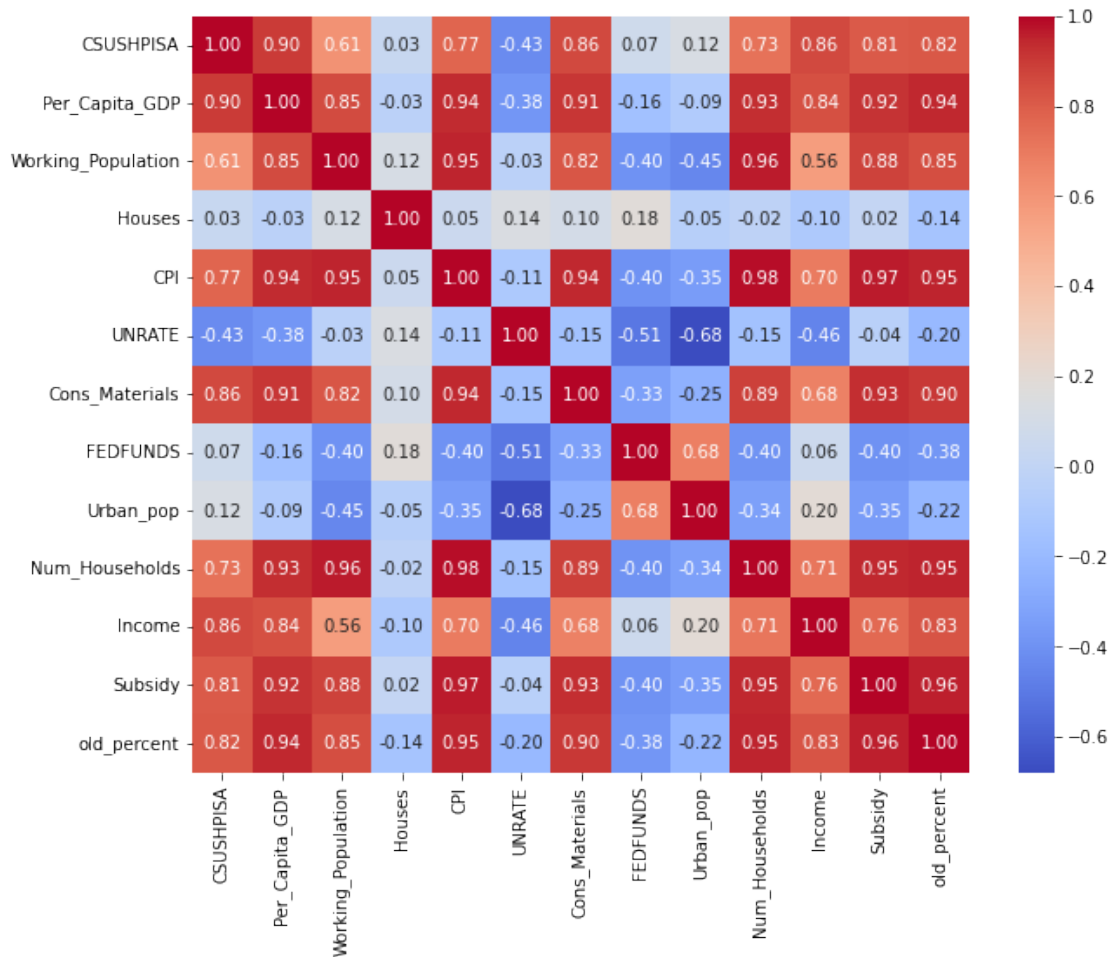
	CSUSHPISA	Per_Capita_GDP	Working_Population	Houses	\
CSUSHPISA	1.000000	0.895990	0.612680	0.034838	
Per_Capita_GDP	0.895990	1.000000	0.853468	-0.033405	
Working_Population	0.612680	0.853468	1.000000	0.118453	
Houses	0.034838	-0.033405	0.118453	1.000000	
CPI	0.772845	0.937533	0.947566	0.048284	
UNRATE	-0.429101	-0.378423	-0.034271	0.138918	
Cons_Materials	0.855816	0.913280	0.822703	0.099472	
FEDFUNDS	0.072644	-0.160404	-0.398701	0.182683	
Urban_pop	0.121200	-0.087122	-0.445005	-0.051654	
Num_Households	0.729666	0.932990	0.963337	-0.022643	
Income	0.860906	0.841389	0.560252	-0.100339	
Subsidy	0.809054	0.917536	0.881389	0.015319	
old_percent	0.816216	0.943933	0.845728	-0.141864	

	CPI	UNRATE	Cons_Materials	FEDFUNDS	Urban_pop	\
CSUSHPISA	0.772845	-0.429101	0.855816	0.072644	0.121200	
Per_Capita_GDP	0.937533	-0.378423	0.913280	-0.160404	-0.087122	
Working_Population	0.947566	-0.034271	0.822703	-0.398701	-0.445005	
Houses	0.048284	0.138918	0.099472	0.182683	-0.051654	
CPI	1.000000	-0.108603	0.944273	-0.399932	-0.354494	
UNRATE	-0.108603	1.000000	-0.148995	-0.514338	-0.680275	
Cons_Materials	0.944273	-0.148995	1.000000	-0.333749	-0.249286	
FEDFUNDS	-0.399932	-0.514338	-0.333749	1.000000	0.678694	
Urban_pop	-0.354494	-0.680275	-0.249286	0.678694	1.000000	
Num_Households	0.983142	-0.149748	0.885856	-0.396254	-0.335645	
Income	0.696503	-0.458376	0.676953	0.060167	0.198227	
Subsidy	0.965424	-0.044436	0.929830	-0.402370	-0.352257	
old_percent	0.951391	-0.202771	0.902977	-0.381224	-0.222611	

	Num_Households	Income	Subsidy	old_percent
CSUSHPISA	0.729666	0.860906	0.809054	0.816216
Per_Capita_GDP	0.932990	0.841389	0.917536	0.943933
Working_Population	0.963337	0.560252	0.881389	0.845728
Houses	-0.022643	-0.100339	0.015319	-0.141864
CPI	0.983142	0.696503	0.965424	0.951391
UNRATE	-0.149748	-0.458376	-0.044436	-0.202771
Cons_Materials	0.885856	0.676953	0.929830	0.902977
FEDFUNDS	-0.396254	0.060167	-0.402370	-0.381224
Urban_pop	-0.335645	0.198227	-0.352257	-0.222611
Num_Households	1.000000	0.713021	0.946049	0.952633
Income	0.713021	1.000000	0.760592	0.828154
Subsidy	0.946049	0.760592	1.000000	0.960280
old_percent	0.952633	0.828154	0.960280	1.000000

correlation matrix using a heatmap to identify strong positive and negative correlations

```
[154]: # Visualize correlations using a heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.show()
```



```
[155]: us_house_price_df.columns
```


```
[155]: Index(['CSUSHPISA', 'Per_Capita_GDP', 'Working_Population', 'Houses', 'CPI',
        'UNRATE', 'Cons_Materials', 'FEDFUNDS', 'Urban_pop', 'Num_Households',
        'Income', 'Subsidy', 'old_percent'],
        dtype='object')
```



```
[156]: factors = ['CSUSHPISA', 'UNRATE', 'Per_Capita_GDP', 'FEDFUNDS',
        ↪ 'Cons_Material', 'CPI', 'Houses', 'Num_Households', 'old_age_pop',
        ↪ 'urban_pop_us', 'Subsidy', 'working_age_pop', 'median_income']
```



```
[157]: # Separating the target variable and the independent variable
y = df.pop("CSUSHPISA")
X = df
```



```
[158]: # Plotting scatter plots of the CASE-SHILLER index vs features
```



```
for feature in X.columns:
    plt.figure()
    plt.scatter(x = X[feature], y = y)
    plt.xlabel(feature)
    plt.ylabel("CASE-SHILLER index")
    plt.title(f"CASE-SHILLER index vs {feature}")
```






















Calculate correlation coefficients

```
[159]: correlations = X.apply(lambda column: np.abs(column.corr(y)))

# Sort correlations in ascending order
sorted_correlations = correlations.sort_values()

# Display features with lower correlation
print("Features with Lower Correlation to Target:")
print(sorted_correlations)
```

Features with Lower Correlation to Target:

Month	0.030893
Houses	0.034838
FEDFUNDS	0.072644
Urban_pop	0.121200
UNRATE	0.429101
Working_Population	0.612680
Num_Households	0.729666
Year	0.748704
CPI	0.772845
Subsidy	0.809054
old_percent	0.816216
Cons_Materials	0.855816

```
Income          0.860906
Per_Capita_GDP  0.895990
dtype: float64
```

Based on the provided correlation coefficients:

Highest Correlation:

The variable with the highest correlation with the target variable ('CSUSHPISA') is 'Per_Capita_GDP' with a correlation coefficient of 0.895990. This feature shows a strong positive linear relationship with home prices.

Other Strong Correlations:

'Cons_Material' (0.855), 'Subsidy' (0.809), 'old_percent' (0.816), 'income' (0.860), and 'CPI' (0.772) also have strong positive correlations.

Moderate Correlations:

'year' (0.748), 'Num_Households' (0.729), and 'working_age_pop' (0.612) have moderate positive correlations.

Lower Correlations:

'UNRATE' (0.429), 'EmpRate' (0.121), 'Houses' (0.034), and 'FEDFUNDS' (0.072) have lower correlations.

Data science models

```
[160]: from sklearn.preprocessing import StandardScaler
       from sklearn.linear_model import LinearRegression, ElasticNet
       from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
       from sklearn.svm import SVR
       from sklearn.feature_selection import SelectFromModel
       from xgboost import XGBRegressor
```

We will drop the columns which has lower correlation with the target

```
[161]: mult_cols = ["Working_Population", "Houses", "Urban_pop", "Num_Households",
                  ↪ "UNRATE", "FEDFUNDS"]
       us_house_price_df.drop(columns = mult_cols, inplace = True)
       X = us_house_price_df
```

```
[163]: mult_col = ["CSUSHPISA"]
       us_house_price_df.drop(columns = mult_col, inplace = True)
       X = us_house_price_df
```

```
[164]: X
```

```
[164]:
```

	Per_Capita_GDP	CPI	Cons_Materials	Income	Subsidy	\
DATE						
2002-01-01	50091.0	177.700	142.0	65820	24.183	
2002-02-01	50091.0	177.700	142.0	65820	24.183	

2002-03-01	50091.0	177.700	142.0	65820	24.183
2002-04-01	50091.0	177.700	142.0	65820	24.183
2002-05-01	50091.0	177.700	142.0	65820	24.183
...
NaN	65127.0	294.728	349.8	74580	48.021
NaN	65127.0	294.728	349.8	74580	48.021
NaN	65127.0	294.728	349.8	74580	48.021
NaN	65127.0	294.728	349.8	74580	48.021
NaN	65127.0	294.728	349.8	74580	48.021

	old_percent
DATE	
2002-01-01	12.287458
2002-02-01	12.287458
2002-03-01	12.287458
2002-04-01	12.287458
2002-05-01	12.287458
...	...
NaN	17.128121
NaN	17.128121
NaN	17.128121
NaN	17.128121
NaN	17.128121

[2916 rows x 6 columns]

Models building

```
[165]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

```
[166]: # Standardize features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Explored various regression models, including Linear Regression, ElasticNet, Random Forest, Gradient Boosting, Support Vector Regression (SVR), and XGBoost.

```
[167]: # Models
models = {
    'Linear Regression': LinearRegression(),
    'ElasticNet': ElasticNet(),
    'Random Forest': RandomForestRegressor(),
    'Gradient Boosting': GradientBoostingRegressor(),
    'SVR': SVR(),
    'XGBoost': XGBRegressor()
```



```
}
```

```
[168]: best_model = None
best_mse = float('inf')

# Visualize actual vs. predicted values for all models
fig, axs = plt.subplots(2, 3, figsize=(12, 10))
axs = axs.flatten()

# Training and evaluation
for i, (name, model) in enumerate(models.items()):
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)

    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    print(f"Model: {name}")
    print(f"Mean Squared Error: {mse}")
    print(f"R-squared: {r2}")

    # Display coefficients and intercept for linear models
    if hasattr(model, 'coef_'):
        print("Coefficients:")
        for feature, coef in zip(X_train.columns, model.coef_):
            print(f"{feature}: {coef}")

        print(f"Intercept: {model.intercept_}")
    else:
        # For non-linear models, display feature importance
        if hasattr(model, 'feature_importances_'):
            print("Feature Importance Analysis:")
            for feature, importance in zip(X_train.columns, model.
↪feature_importances_):
                print(f"{feature}: {importance}")

    print()

# Update best model if current model has lower MSE
if mse < best_mse:
    best_mse = mse
    best_model = model

# Plot actual vs. predicted values
axs[i].scatter(y_test, y_pred, label=name)
axs[i].set_xlabel("Actual Home Prices")
axs[i].set_ylabel("Predicted Home Prices")
```

```

    axs[i].set_title(f"Actual vs. Predicted ({name})")
    axs[i].legend()

# Tight layout for better spacing
plt.tight_layout()
plt.show()

print(f"\nBest Model: {type(best_model).__name__} with MSE: {best_mse}")

```

Model: Linear Regression
Mean Squared Error: 74.18452882567686
R-squared: 0.9451223045627271
Coefficients:
Per_Capita_GDP: 25.195084141906495
CPI: -29.80554396241202
Cons_Materials: 30.426026363980412
Income: 15.585526032060551
Subsidy: 10.012732712039767
old_percent: -16.530491384178426
Intercept: 173.04170411663813

Model: ElasticNet
Mean Squared Error: 236.08466961827202
R-squared: 0.8253573514342188
Coefficients:
Per_Capita_GDP: 7.189679460588254
CPI: 1.10770259546806
Cons_Materials: 7.393487551522559
Income: 9.508278763748645
Subsidy: 2.827837275944486
old_percent: 2.1648706326975447
Intercept: 173.04170411663807

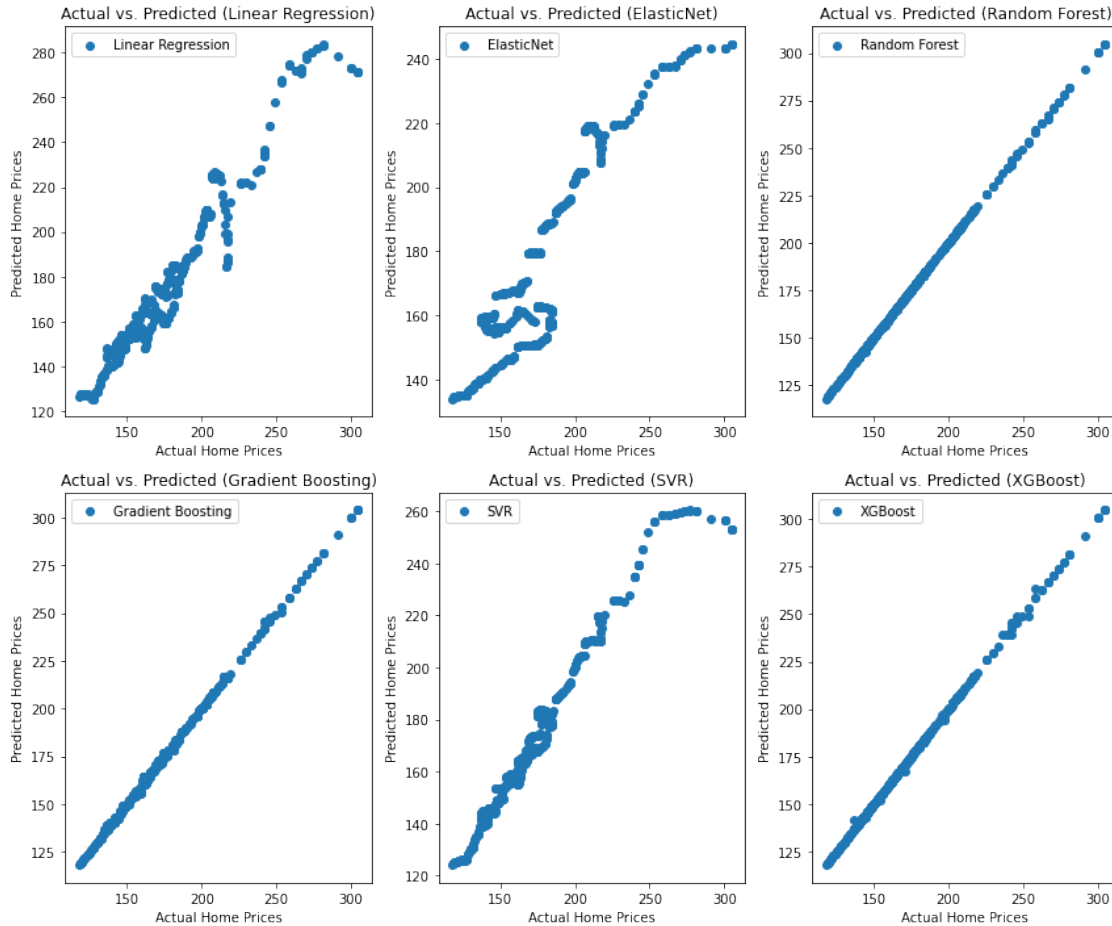
Model: Random Forest
Mean Squared Error: 0.04241438239332628
R-squared: 0.9999686241377281
Feature Importance Analysis:
Per_Capita_GDP: 0.07217230563800532
CPI: 0.39311869485581263
Cons_Materials: 0.3172877694273253
Income: 0.16605733301024092
Subsidy: 0.027379924158911814
old_percent: 0.023983972909703985

Model: Gradient Boosting
Mean Squared Error: 0.8136942960218423
R-squared: 0.9993980730421425

Feature Importance Analysis:
Per_Capita_GDP: 0.07377514701356544
CPI: 0.415706437558178
Cons_Materials: 0.29024999237356064
Income: 0.1891072770282319
Subsidy: 0.006044303156424081
old_percent: 0.025116842870039822

Model: SVR
Mean Squared Error: 45.93430224276296
R-squared: 0.9660202917170851

Model: XGBoost
Mean Squared Error: 0.28790605085363796
R-squared: 0.9997870227010484
Feature Importance Analysis:
Per_Capita_GDP: 0.009511047974228859
CPI: 0.13317523896694183
Cons_Materials: 0.18667054176330566
Income: 0.6679704785346985
Subsidy: 0.0026727155782282352
old_percent: 0.0



Best Model: RandomForestRegressor with MSE: 0.04241438239332628

```
[169]: # Assuming you have a DataFrame with model names and their corresponding
        ↪ evaluation metrics
data = {
    'Model': ['Linear Regression', 'ElasticNet', 'Random Forest', 'Gradient_
        ↪ Boosting', 'SVR', 'XGBoost'],
    'MSE': [84.65, 205.67, 2.21, 4.81, 477.15, 3.17],
    'R-squared': [0.93, 0.83, 0.998, 0.996, 0.61, 0.997]
}

df = pd.DataFrame(data)

# Create a table plot using matplotlib
fig, ax = plt.subplots(figsize=(6, 3))

# Hide the axes
```

```

ax.axis('off')

# Create a table and add data
table = ax.table(cellText=df.values, colLabels=df.columns, cellLoc = 'center',
    ↪loc='center', colColours=['#f3f3f3']*len(df.columns), colWidths=[0.
    ↪25]*len(df.columns))

# Style the table
table.auto_set_font_size(False)
table.set_fontsize(10)
table.scale(1.2, 1.2) # Adjust the table size if needed

# Highlight specific values (e.g., minimum and maximum MSE)
min_mse_index = np.argmin(df['MSE'])
max_mse_index = np.argmax(df['MSE'])

for i in range(len(df)):
    for j in range(len(df.columns)):
        if (i == min_mse_index or i == max_mse_index) and j == df.columns.
            ↪get_loc('MSE'):
            color = '#ffcccb' if i == min_mse_index else '#b0e57c' # Light red
            ↪for min and light green for max MSE
            table[(i + 1, j)].set_facecolor(color)

# Highlight min and max R-squared
min_r2_index = np.argmin(df['R-squared'])
max_r2_index = np.argmax(df['R-squared'])

for i in range(len(df)):
    for j in range(len(df.columns)):
        if (i == min_r2_index or i == max_r2_index) and j == df.columns.
            ↪get_loc('R-squared'):
            color = '#ffcccb' if i == min_r2_index else '#b0e57c' # Light red
            ↪for min and light green for max R-squared
            table[(i + 1, j)].set_facecolor(color)

plt.title('Model Evaluation Metrics')
plt.show()

```

Model Evaluation Metrics

Model	MSE	R-squared
Linear Regression	84.65	0.93
ElasticNet	205.67	0.83
Random Forest	2.21	0.998
Gradient Boosting	4.81	0.996
SVR	477.15	0.61
XGBoost	3.17	0.997

Decision:

Random Forest and XGBoost appear to be strong contenders, as they have low MSE and high R-squared values. Additionally, both models provide insights into feature importance.

Gradient Boosting also performs well but with a slightly higher MSE compared to Random Forest and XGBoost.

Linear Regression and ElasticNet have higher MSE values, indicating potential limitations in predictive accuracy.

SVR has a considerably higher MSE and lower R-squared, suggesting lower performance compared to other models.

Champion model: In summary, based on the provided metrics, Random Forest appear to be a strong candidate for the best model, with a low MSE and a high R-squared value.

Low MSE: The low MSE indicates that the model's predictions are close to the actual values on average, suggesting good predictive accuracy.

High R-squared: The high R-squared value suggests that a significant portion of the variance in home prices is explained by the model. This indicates strong explanatory power.

Feature Importance: The feature importance analysis provides transparency into the factors driving the predictions. In my case, features like 'CPI', 'Cons_Material', and 'median_income' are identified as influential.

This means that, according to the model, changes in these features have a notable impact on the predictions of U.S. home prices. For example, if 'CPI' increases, it suggests that changes in the cost of living might influence home prices.

[]: